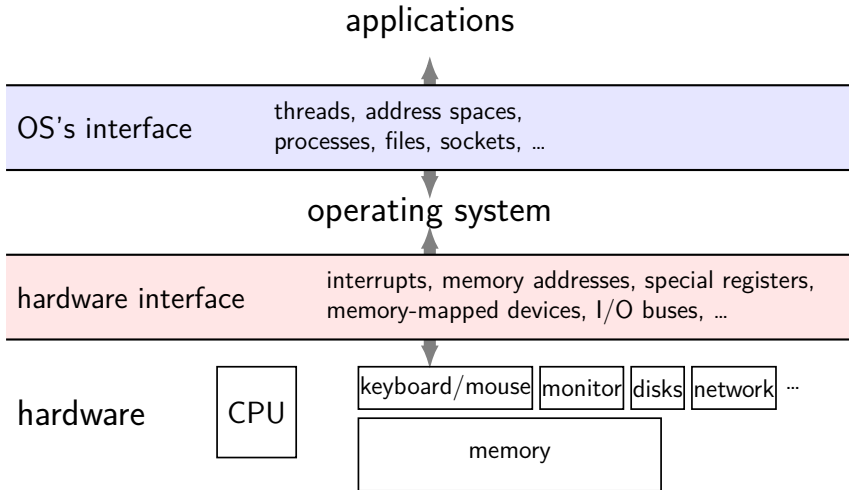


the abstract virtual machine



abstract VM: application view

applications



OS's interface


threads, address spaces,
processes, files, sockets, ...

the application's “machine” **is the operating system**

no hardware I/O details visible — future-proof

more featureful interfaces than real hardware

abstract VM: OS view



operating system's job: translate one interface to another

program → process → CPU and memory

application 1

applications

OS's interface

threads, address spaces,
processes, files, sockets, ...

operating system

hardware interface

interrupts, memory addresses, special registers,
memory-mapped devices, I/O buses, ...

hardware

CPU

keyboard/mouse

monitor

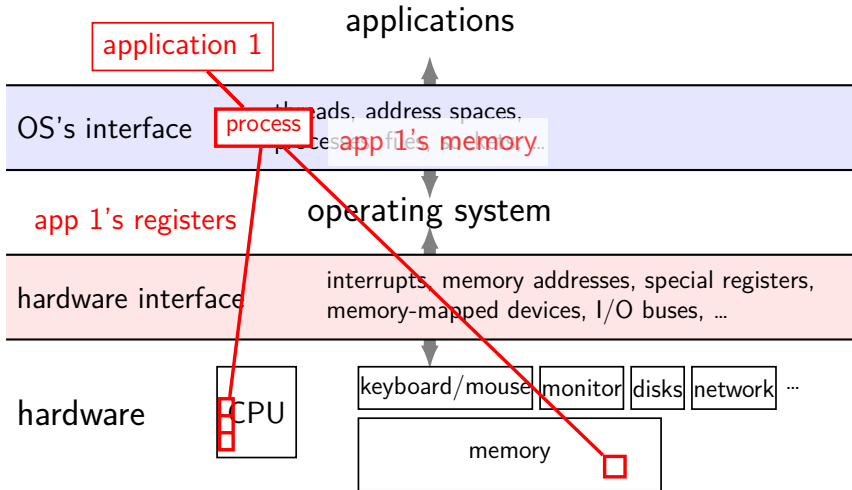
disks

network

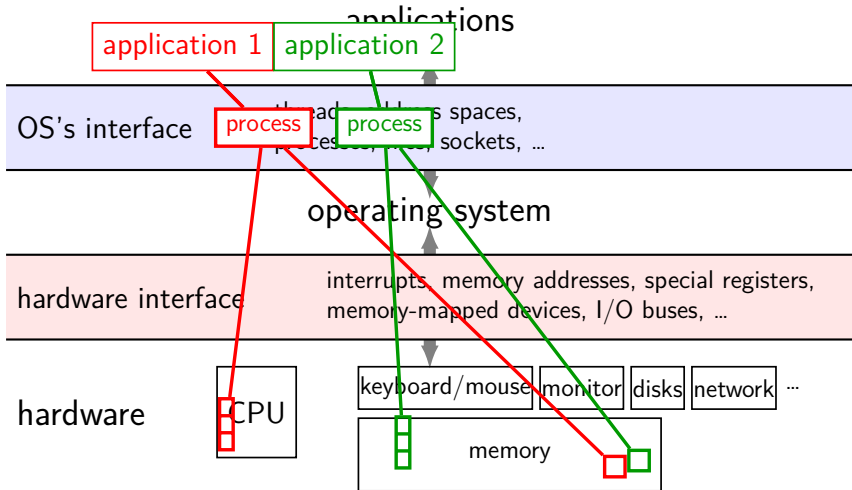
...

memory

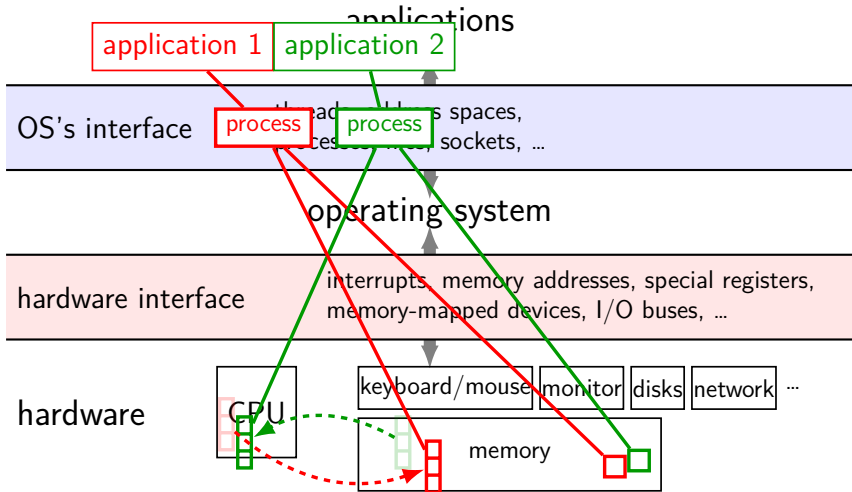
program → process → CPU and memory



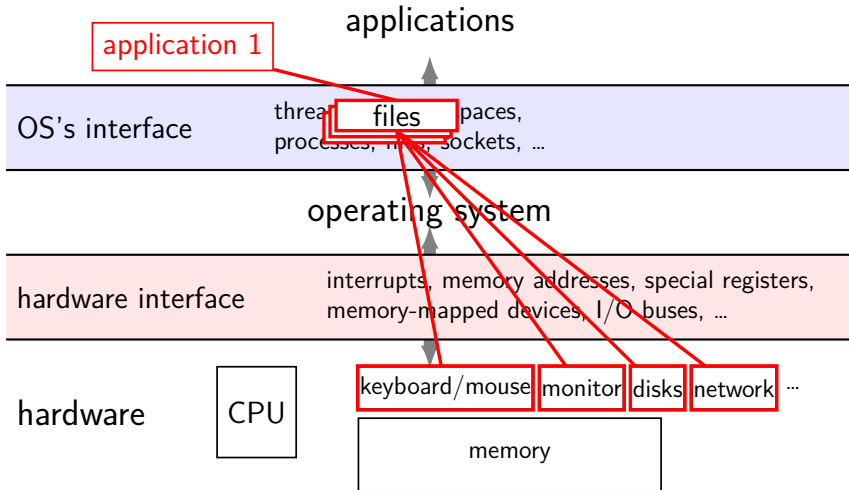
program → process → CPU and memory

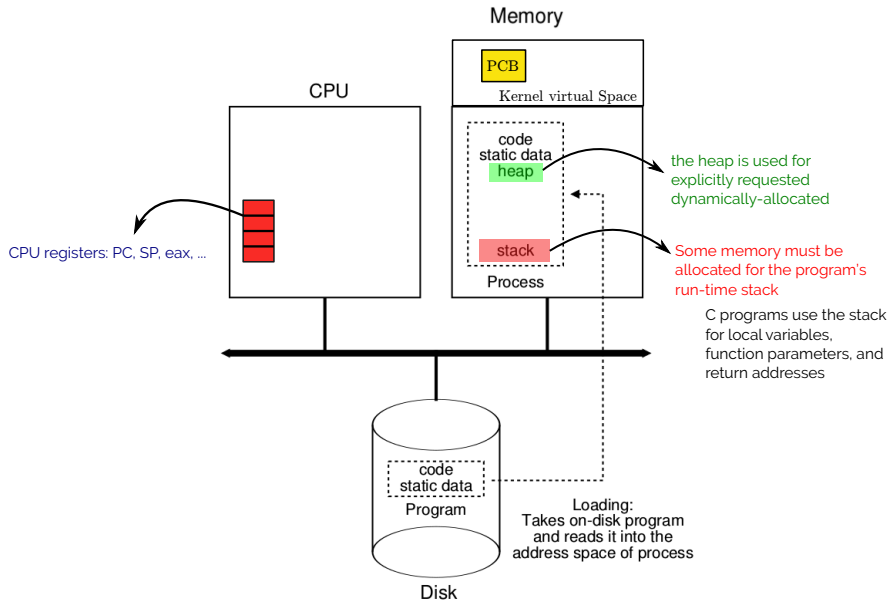



program → process → CPU and memory



files → input/output







- running: its state word is contained in a processor which is running.
- ready: it could be placed in execution by a processor if one were free
- suspend: awaiting activation by an external event, such as the completion of an i/o function.

Example

Time	Process ₀	Process ₁	Notes
1	Running	Ready	
2	Running	Ready	
3	Running	Ready	
4	Running	Ready	Process ₀ now done
5	–	Running	
6	–	Running	
7	–	Running	
8	–	Running	Process ₁ now done

Tracing Process State: CPU Only

Example

Time	Process ₀	Process ₁	Notes
1	Running	Ready	
2	Running	Ready	
3	Running	Ready	Process ₀ initiates I/O
4	Blocked	Running	Process ₀ is blocked,
5	Blocked	Running	so Process ₁ runs
6	Blocked	Running	
7	Ready	Running	I/O done
8	Ready	Running	Process ₁ now done
9	Running	–	
10	Running	–	Process ₀ now done

Tracing Process State: CPU and I/O

ASIDE: xv6

In this course, instead of practical OS kernels (with millions of code lines), we use xv6 (just few thousands of code lines) for Kernel hacking

- A teaching OS developed at MIT
- simplified Unix version 6 (main change: from PDP-11 to Intel x86)
- well structured and documented
- github link: <https://github.com/mit-pdos/xv6-public>
- A commentary book:
<https://pdos.csail.mit.edu/6.828/2014/xv6/book-rev8.pdf>
- visit <https://pdos.csail.mit.edu/6.828/2012/xv6.html>

```

// the registers xv6 will save and restore
// to stop and subsequently restart a process
struct context {
    int eip;
    int esp;
    int ebx;
    int ecx;
    int edx;
    int esi;
    int edi;
    int ebp;
};

// the different states a process can be in
enum proc_state { UNUSED, EMBRYO, SLEEPING,
    RUNNABLE, RUNNING, ZOMBIE };

// the information xv6 tracks about each process
// including its register context and state
struct proc {
    char *mem;                // Start of process memory
    uint sz;                  // Size of process memory
    char *kstack;             // Bottom of kernel stack
                                // for this process
    enum proc_state state;    // Process state
    int pid;                  // Process ID
    struct proc *parent;      // Parent process
    void *chan;               // If !zero, sleeping on chan
    int killed;               // If !zero, has been killed
    struct file *ofile[NOFILE]; // Open files
    struct inode *cwd;        // Current directory
    struct context context;    // Switch here to run process
    struct trapframe *tf;     // Trap frame for the
                                // current interrupt
};

```

The xv6 Proc Structure (Process Control Block (PCB)/process descriptor)

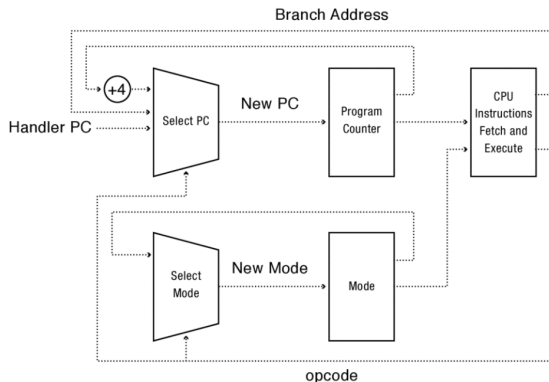
Hardware Support

Dual Mode

Exception

Memory Protection

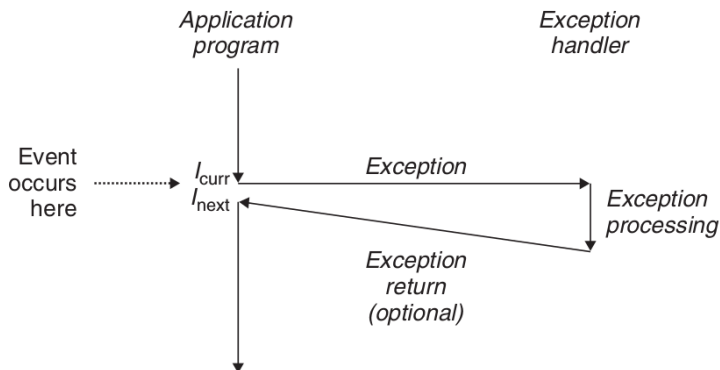
The operation of a CPU with kernel and user modes


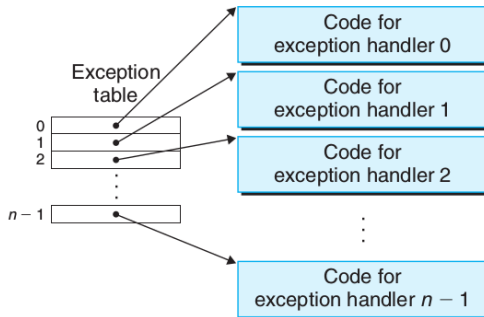


The kernel/user mode bit is one flag in the processor **status register**, set whenever the kernel is entered and reset whenever the kernel switches back to user mode.

Exception

An exception is an abrupt change in the control flow in response to some change in the processor's state.





Classes of Exceptions

Asynchronous (Interrupts)

Synchronous: Traps, Faults, Aborts

Exercise: How many exceptions?

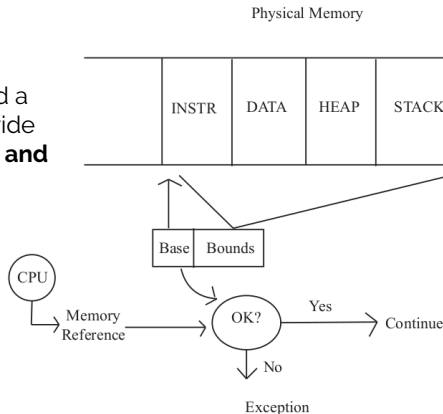
Single core, 3 processes A, B, and C

- Running process A
- A prompts for input then
- A waits to read a keypress
- While A is waiting for the keypress, OS runs B, then C
- then keypress happens, and OS switches to A immediately
- then A exits

Memory Protection

All memory accesses outside of a process's valid memory region are prohibited when executing in user mode.

Early computers pioneered a simple mechanism to provide protection, known as **base and bounds**.



More elaborated methods in the next chapter